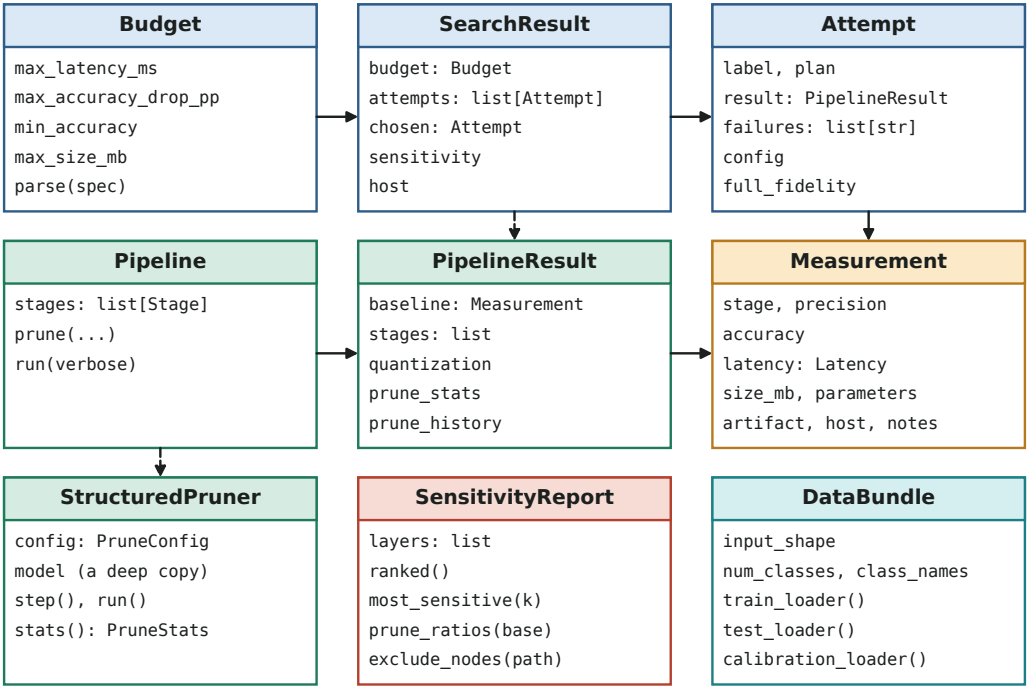
Budget
max_latency_ms
max_accuracy_drop_pp
min_accuracy
max_size_mb
parse(spec)
SearchResult
budget: Budget
attempts: list[Attempt]
chosen: Attempt
sensitivity
host
Attempt
label, plan
result: PipelineResult
failures: list[str]
config
full_fidelity
Pipeline
stages: list[Stage]
prune(...)
run(verbose)
PipelineResult
baseline: Measurement
stages: list
quantization
prune_stats
prune_history
Measurement
stage, precision
accuracy
latency: Latency
size_mb, parameters
artifact, host, notes
StructuredPruner
config: PruneConfig
model (a deep copy)
step(), run()
stats(): PruneStats
SensitivityReport
layers: list
ranked()
most_sensitive(k)
prune_ratios(base)
exclude_nodes(path)
DataBundle
input_shape
num_classes, class_names
train_loader()
test_loader()
calibration_loader()