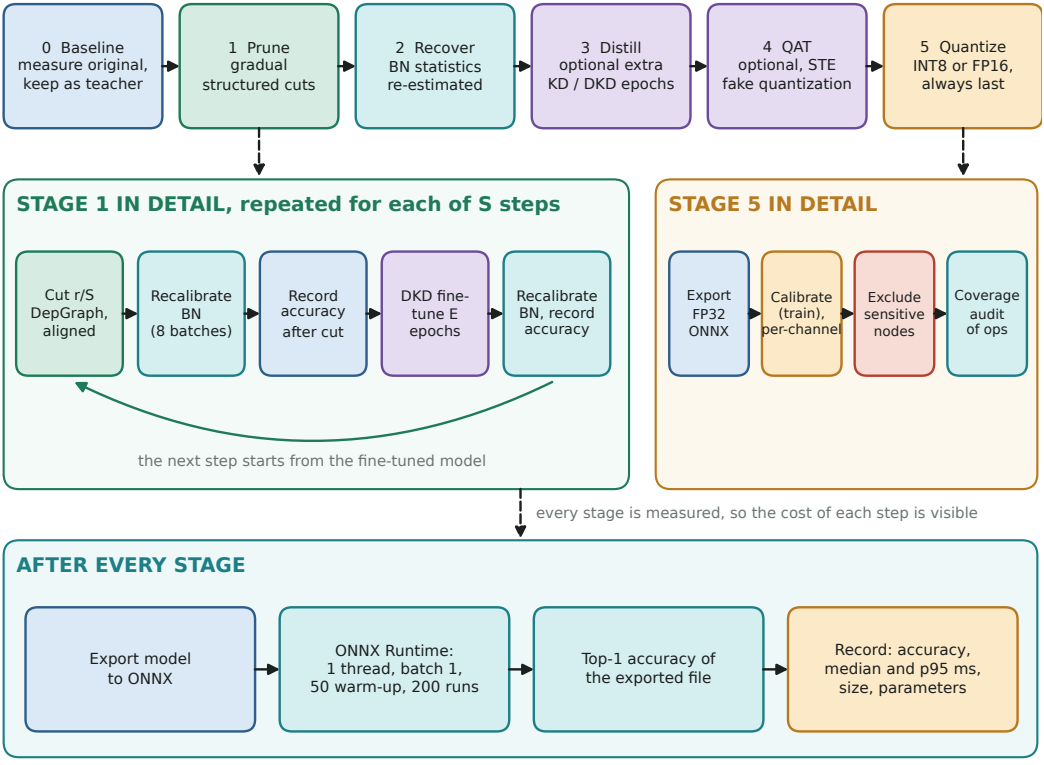
0 Baseline measure original, keep as teacher
1 Prune gradual structured cuts
2 Recover BN statistics re-estimated
3 Distill optional extra KD / DKD epochs
4 QAT optional, STE fake quantization
5 Quantize INT8 or FP16, always last
STAGE 1 IN DETAIL, repeated for each of S steps
Cut r/S DepGraph, aligned
Recalibrate BN (8 batches)
Record accuracy after cut
DKD fine-tune E epochs
Recalibrate BN, record accuracy
the next step starts from the fine-tuned model
STAGE 5 IN DETAIL
Export FP32 ONNX
Calibrate (train), per-channel
Exclude sensitive nodes
Coverage audit of ops
every stage is measured, so the cost of each step is visible
AFTER EVERY STAGE
Export model to ONNX
ONNX Runtime: 1 thread, batch 1, 50 warm-up, 200 runs
Top-1 accuracy of the exported file
Record: accuracy, median and p95 ms, size, parameters