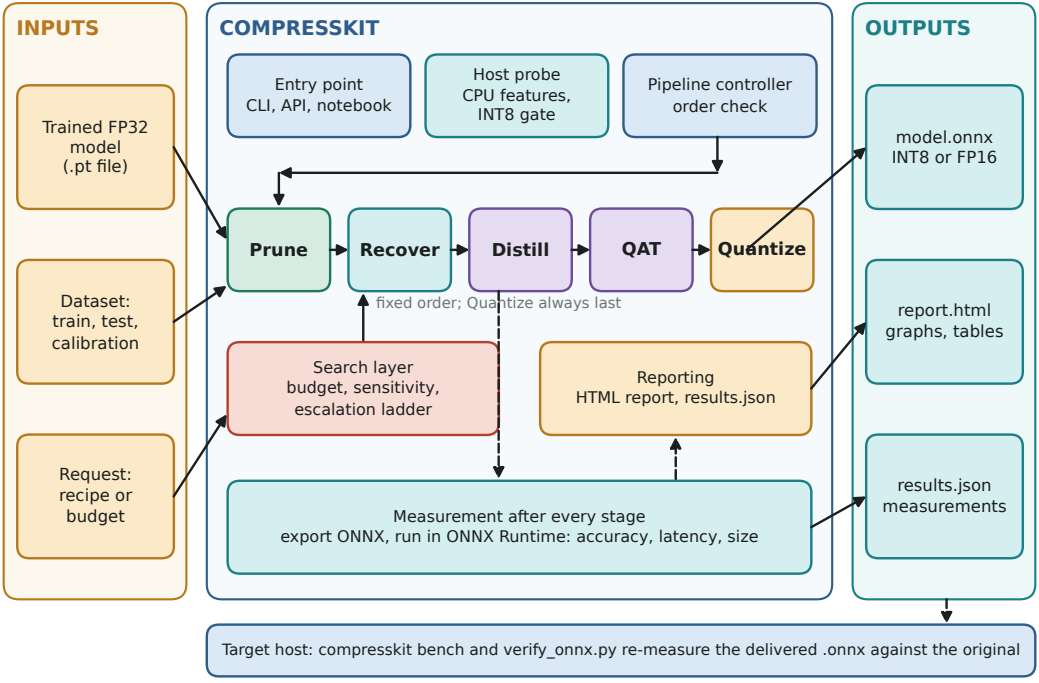
INPUTS
Trained FP32 model (.pt file)
Dataset: train, test, calibration
Request: recipe or budget
COMPRESSKIT
Entry point CLI, API, notebook
Host probe CPU features, INT8 gate
Pipeline controller order check
Prune
Recover
Distill
QAT
Quantize
fixed order; Quantize always last
Search layer budget, sensitivity, escalation ladder
Reporting HTML report, results.json
Measurement after every stage export ONNX, run in ONNX Runtime: accuracy, latency, size
OUTPUTS
model.onnx INT8 or FP16
report.html graphs, tables
results.json measurements
Target host: compresskit bench and verify_onnx.py re-measure the delivered .onnx against the original